





MoE variants vs baselines (7.86B tokens)
x-axis: total params



† BoltzmannMoE uses soft routing: all K experts compute every step (active = total params). TopK uses hard routing: only top- k experts active per token.
Dashed line: same H-series arch, different routing (h1_topk $\leftrightarrow$ h1_boltz_fullsize).